

# Operating Systems

## Assignment-2

Rohan Jain

2019095

### Question 1

Yes, there is a difference between the value of the global variable printed in the two cases. In case 1, `fork()` duplicates the process completely using COW principles. However, `pthread_create()` only copies the **code** and **stack** section of the process. This means that both child and parent thread point to the same global variable. In `fork()`, the parent and child processes have their own copies of the global variable.

In the case of `fork()`, if the parent process has counted up to 35 and then context switch happens, the child doesn't start from 35, it starts from 10. Suppose the child now counts down to -1 and context switch happens, the parent will start from its value, i.e. 35 and not -1. This is because both processes have their own memory allocated for that global variable.

In the case of `pthread_create()`, if the parent process has counted up to 35 and then context switch happens, the child starts from 35 and not 10. Suppose the child now counts down to -1 and context switch happens, the parent now starts from -1 and not 35. This is because both processes have only one shared memory for that global variable in the common **data** section.

#### `fork()` sample:

```
Parent Process: 40
Parent Process: 41
Parent Process: 42
Parent Process: 43
Parent Process: 44
Parent Process: 45
Parent Process: 46
Parent Process: 47
Parent Process: 48
Child process running! PID is: 40364
Parent Process: 49
Parent Process: 50
Child Process: 9
Parent Process: 51
Child Process: 8
Parent Process: 52
Child Process: 7
Parent Process: 53
Child Process: 6
Parent Process: 54
Child Process: 5
Parent Process: 55
Child Process: 4
```

### pthread\_create() sample:

```
Parent Thread running!  
Parent Thread: 11  
Parent Thread: 12  
Parent Thread: 13  
Parent Thread: 14  
Parent Thread: 15  
Parent Thread: 16  
Parent Thread: 17  
Parent Thread: 18  
Parent Thread: 19  
Parent Thread: 20  
Child Thread created!  
Child Thread: 20  
Child Thread: 19  
Child Thread: 18  
Child Thread: 17  
Child Thread: 16  
Child Thread: 15  
Child Thread: 14  
Child Thread: 13  
Parent Thread: 21  
Parent Thread: 13  
Parent Thread: 14  
Parent Thread: 15  
Parent Thread: 16  
Parent Thread: 17  
Parent Thread: 18  
Parent Thread: 19  
Parent Thread: 20  
Parent Thread: 21  
Parent Thread: 22  
Child Thread: 12  
Child Thread: 22  
Child Thread: 21  
Child Thread: 20
```

In both cases, the parent thread/process eventually reaches `100` and child thread/process reaches `-90`. However, in case of thread there is a possibility that the program never ends. If the parent increments the number 10 times, context switch happens, and then child process decrements number 10 times and again context switch happens. In this case, the process may never end. The chances of this happening are very less though and the process finishes almost every time with varying time taken.

## Question 2

### Description

The system call prints PID, State, Real time priority, Scheduling policy, Number of CPUs allowed, Command name and the number of context switches done of the process corresponding to `pid` on kernel log and in a file given by the parameter `filename` in the current working directory.

### Function definition

```
long sys_sh_task_info(pid_t pid, char* filename);
```

Returns 0 on success and negative number on error. It sets `errno` on return.

To check the output of the syscall on kernel log: `dmesg | tail`

### Explanation of Code:

```
SYSCALL_DEFINE2(sh_task_info, pid_t, pid, char *, file_name)
```

This is the main system call function that is defined using the `SYSCALL_DEFINE` macro. To find the `task_struct` corresponding to the `pid` passed in parameter I have used functions `find_vpid()` and `pid_task()`.

```
task = pid_task(find_vpid(pid), PIDTYPE_PID);
```

If the `task_struct` corresponding to the `pid` is not found, the `errno` is set to `ESRCH` = No such process and the returned value is `-ESRCH`.

Next, the `set_output()` function is called.

```
int set_output(char *output, struct task_struct *task)
```

The function `set_output` populates the string output with the fields given in description. It takes the fields from the `task_struct` pointed by `task`. It makes use of functions like `snprintf()` and `strncat()` which are defined in the kernel API.

To access the string passed in filename, I have to copy the string from user space to kernel space. To do this, I have used the function `copy_from_user()`

```
copy_from_user(filepath, file_name, 256);
```

If the `filename` given is greater than 256 bytes, the `errno` is set to `ENAMETOOLONG` = File name too long and the returned value is `-ENAMETOOLONG`.

If the `filename` given is invalid, then `errno` is set to `ENOENT` = No such file or directory and the returned value is `-ENOENT`.

I open/create the file using `filp_open()`. I give the permissions `0666` and set the flags to write and create.

```
fileptr = filp_open(filepath, O_WRONLY | O_CREAT, 0666);
```

After getting the `fileptr` (`struct file*`) , I write the previously generated output string to the file pointed by `fileptr` using `kernel_write()`.

```
returnVal = kernel_write(fileptr, output, len, &pos);
```

If `kernel_write()` does not write the whole message, then `errno` is set to `EIO` = I/O Error and the returned value is `-EIO`.

I close the file using `filp_close()` and return 0 as everything worked if we have reached this command.

```
filp_close(fileptr, NULL);
```

## Expected Input Output

The user has to call the system call and give 2 inputs `pid_t pid` and `char *filename`. `pid` refers to the pid of the process whose details have to be printed and `filename` refers to the name of the file in which the details are stored.

The output of the system call is all the details of the values mentioned in description. The string representations of the numeric values like scheduling policy are also mentioned in the printed values.

## All errors returned by code:

ESRCH  
ENAMETOOLONG  
ENOENT  
EIO

The conditions for their return are explained above.

## Test Cases

### Sample Output 1:

```
rohanj-02@ubuntu: ~  
rohanj-02@ubuntu:~$ gcc -o test test.c  
rohanj-02@ubuntu:~$ ./test  
Enter pid(-1 to send pid of this process): 1  
Enter filename: file1.txt  
Run dmesg or see the file in your current directory to check the output.  
rohanj-02@ubuntu:~$ cat file1.txt  
PID: 1  
State: 1 (TASK_INTERRUPTIBLE)  
Real Time Priority: 0  
Scheduling Policy: 0 (SCHED_NORMAL)  
Number of CPUs allowed: 128  
Command name: systemd  
Number of context switch done: 3500  
rohanj-02@ubuntu:~$ dmesg | tail  
[ 1032.604895] hrtimer: interrupt took 11349938 ns  
[ 5404.036246] Running sh_task_info syscall  
[ 5404.036271] PID: 1  
State: 1 (TASK_INTERRUPTIBLE)  
Real Time Priority: 0  
Scheduling Policy: 0 (SCHED_NORMAL)  
Number of CPUs allowed: 128  
Command name: systemd  
Number of context switch done: 3500  
rohanj-02@ubuntu:~$
```

## Sample Output 2:

```
rohanj-02@ubuntu:~$ ./test  
Enter pid(-1 to send pid of this process): 26  
Enter filename: file2.txt  
Run dmesg or see the file in your current directory to check the output.  
rohanj-02@ubuntu:~$ cat file2.txt  
PID: 26  
State: 1026 (Unknown)  
Real Time Priority: 0  
Scheduling Policy: 0 (SCHED_NORMAL)  
Number of CPUs allowed: 1  
Command name: kworker/2:0H  
Number of context switch done: 8  
rohanj-02@ubuntu:~$ dmesg | tail  
[ 5555.544712] e1000: ens33 NIC Link is Up 1000 Mbps Full Duplex, Flow Control: None  
[ 5569.527994] Running sh_task_info syscall  
[ 5569.528015] PID: 26  
State: 1026 (Unknown)  
Real Time Priority: 0  
Scheduling Policy: 0 (SCHED_NORMAL)  
Number of CPUs allowed: 1  
Command name: kworker/2:0H  
Number of context switch done: 8
```

## Error Handling

To test ESRCH, send an invalid pid like a negative number or a really large number.

To test ENOENT, send a filename such as hello/hello.txt where hello/ directory does not exist.

```
rohanj-02@ubuntu:~$ ./test
Enter pid(-1 to send pid of this process): 3146781
Enter filename: hello.txt
Error: sh_task_info: No such process
rohanj-02@ubuntu:~$ ./test
Enter pid(-1 to send pid of this process): -1
Enter filename: hel/hello.txt
Error: sh_task_info: No such file or directory
rohanj-02@ubuntu:~$ ls
Desktop  Downloads  file2.txt  init.txt  Pictures  snap      test      Videos
Documents  file1.txt  file.txt  Music     Public    Templates test.c
```

## Running test.c

To ensure giving a valid pid and checking the values, run `top` in another terminal tab and compare the values from there.